

# SMAI Assignment 5

## KDE

### KDE Class

```
class KDE:  
    def __init__(self, bandwidth=1.0, kernel='gaussian'):  
        self.bandwidth = bandwidth  
        self.kernel = kernel  
        self.data = None  
  
    def _box_kernel(self, u):  
        """Box kernel: Returns 0.5 if |u| <= 1, else 0."""  
        return np.where(np.abs(u) <= 1, 0.5, 0)  
  
    def _gaussian_kernel(self, u):  
        """Gaussian kernel."""  
        return (1 / np.sqrt(2 * np.pi)) * np.exp(-0.5 * u ** 2)  
  
    def _triangular_kernel(self, u):  
        """Triangular kernel: Returns 1 - |u| if |u| <= 1, else 0."""  
        return np.where(np.abs(u) <= 1, 1 - np.abs(u), 0)  
  
    def _select_kernel(self, u):  
        """Selects the kernel function based on the specified kernel type.  
        if self.kernel == 'box':  
            return self._box_kernel(u)  
        elif self.kernel == 'gaussian':  
            return self._gaussian_kernel(u)  
        elif self.kernel == 'triangular':  
            return self._triangular_kernel(u)  
        else:  
            raise ValueError("Unsupported kernel type. Choose 'box', 'gaussian' or 'triangular'.")
```

```


def fit(self, data):
    """Fit the KDE model to data."""
    self.data = data

def predict(self, x):
    """Calculate the density estimation at a single point
    distances = (x - self.data) / self.bandwidth
    kernel_values = self._select_kernel(distances)
    density = np.sum(kernel_values, axis=0) / (len(self.data))
    return np.sum(density) # Ensure a single density value


def score_samples(self, points):
    """Calculate the density estimation for multiple points
    densities = np.array([self.predict(point) for point in points])
    return np.log(densities) # Return log density for consistency


```

## Generate synthetic data




## KDE vs GMM






The KDE model estimates the density better.


GMM estimation improves as the number of components are increased





# HMM

## Dataset



## Model Architecture

```
# Initialize HMM models and train them
def train_hmm_models(train_data):
    models = []
    for digit, features in train_data.items():
        model = hmm.GaussianHMM(n_components=4, covariance_type='diag')
        X = np.vstack(features) # Stack all MFCC features for
                                # this digit into one matrix
```

```
    lengths = [len(f) for f in features] # Lengths of each feature
    model.fit(X, lengths) # Train the HMM model
    models[digit] = model
return models
```

```
# Predict the digit based on highest HMM model score
def predict_digit(mfcc_features, models):
    max_score = float("-inf")
    best_digit = None
    for digit, model in models.items():
        score = model.score(mfcc_features)
        if score > max_score:
            max_score = score
            best_digit = digit
    return best_digit
```

## Metrics

On trained data

```
Test Accuracy: 89.00%
Validation Accuracy: 91.33%
Train Accuracy: 92.71%
```

On own audio

```
File: _2.wav, Predicted Digit: 2
File: _8.wav, Predicted Digit: 8
File: _4.wav, Predicted Digit: 4
File: _3.wav, Predicted Digit: 9
File: _0.wav, Predicted Digit: 1
File: _1.wav, Predicted Digit: 0
File: _6.wav, Predicted Digit: 6
File: _5.wav, Predicted Digit: 5
```

```
File: _7.wav, Predicted Digit: 7
File: _9.wav, Predicted Digit: 9
```

Accuracy = 70 percent

# RNN

## Counting Bits

### Dataset

Training examples:

		sequence	count_of_ones
75721	[0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1]		7
80184	[0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0]		5
19864		[1, 0, 0, 1]	2
76699		[1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0]	5
92991	[1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1]		5

Validation examples:

		sequence	count_of_ones
52882		[0, 0, 0, 1, 1, 0, 1]	3
9835		[0, 1, 1, 1, 0, 0, 0, 0]	3
42963	[0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0]		5
5380		[0, 0, 1, 1, 1]	3
69603	[1, 1, 0, 1, 0, 0, 0, 0, 1, 0]		4

Test examples:

		sequence	count_of_ones
5		[0, 0, 1, 1, 1]	3
15	[0, 1, 0, 1, 0, 1, 0]		3
20		[0]	0
37		[0, 1, 1, 1]	3
42	[0, 1, 0, 0, 1, 0, 0]		2

### Architecture

```
# Define the RNN model
class RNNBinaryCounter(nn.Module):
    def __init__(self, input_size=1, hidden_size=16, output_s
        super(RNNBinaryCounter, self).__init__()
        self.hidden_size = hidden_size
```

```

        self.rnn = nn.RNN(input_size, hidden_size, batch_firs
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        # Initialize hidden state with zeros
        h0 = torch.zeros(1, x.size(0), self.hidden_size).to(x

            # Forward propagate through RNN
            out, _ = self.rnn(x.unsqueeze(-1), h0)

            # Use the last output from the RNN for the final coun
            out = self.fc(out[:, -1, :])
            return out

```

## Training

```

# Initialize lists to store MAE values for each epoch
train_mae_history = []
val_mae_history = []

num_epochs = 10
for epoch in range(num_epochs):
    model.train()
    train_loss = 0
    train_mae = 0

    for sequences, labels in train_loader:
        sequences, labels = sequences.to(device), labels.to(d

            # Forward pass
            outputs = model(sequences)
            loss = criterion(outputs.squeeze(), labels) # MSE lo

            # Backward and optimize
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

```

```

    train_loss += loss.item()

    # Calculate MAE manually for training
    mae = torch.abs(outputs.squeeze() - labels).mean().item()
    train_mae += mae
    train_loss /= len(train_loader)
    train_mae /= len(train_loader)
    train_mae_history.append(train_mae)


    # Validation - Calculate MAE while keeping MSE as the loss
    model.eval()
    val_loss = 0
    val_mae = 0
    with torch.no_grad():
        for sequences, labels in val_loader:
            sequences, labels = sequences.to(device), labels.to(device)
            outputs = model(sequences)

            # Calculate MSE Loss
            loss = criterion(outputs.squeeze(), labels)
            val_loss += loss.item()

            # Calculate MAE manually for validation
            mae = torch.abs(outputs.squeeze() - labels).mean().item()
            val_mae += mae
    val_loss /= len(val_loader)
    val_mae /= len(val_loader)
    val_mae_history.append(val_mae)

    print(f"Epoch [{epoch+1}/{num_epochs}], Train Loss (MSE): {train_loss:.4f}, Train MAE: {train_mae:.4f}, Val MAE: {val_mae:.4f}")

```




Model Test MAE: 0.0269

Random Baseline Test MAE: 5.4843

The model performs better than the random baseline.

## Generalization



### Trend:

Since the model was trained on sequence lengths between 1 and 16, it is expected to perform better on sequences of these lengths, resulting in relatively low MAE values. For sequence lengths beyond this range (17–32), the model is likely to experience a higher MAE as it encounters lengths it has never seen during training.

### Interpretation:

Sequence Lengths 1–16: The MAE is lower and stable as the model has learned to count within this range during training. Sequence Lengths 17–32: A progressive increase in MAE is expected as the sequence length increases, as the model has to generalize counting accurately beyond its training distribution.

## Optical Character Recognition

### Dataset

dispirit.png	freckly.png	hoer.png	ascertainer.png	bumpily.png
dispirit	freckly	hoer	ascertainer	bumpily

With encoded labels

Label: tensor([ 9, 14, 20, 5, 18, 3, 1, 18, 4, 9, 14, 1, 12, 0, 0, 0])

intercardinal

Label: tensor([ 3, 25, 3, 12, 9, 26, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

cyclize

Label: tensor([ 7, 21, 26, 26, 12, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

guzzle

## Architecture

```
import torch.nn.functional as F

# Sample CNN-RNN model
class CNNEncoder(nn.Module):
    def __init__(self):
        super(CNNEncoder, self).__init__()
        self.conv1 = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        # print(x.shape)
        x = self.pool(F.relu(self.conv2(x)))
        # print(x.shape)
```

```

        # return x.permute(0, 2, 3, 1) # Batch x Width x H
eight x Channels
        return x

class RNND decoder(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim):
        super(RNND decoder, self).__init__()
        self.rnn = nn.LSTM(input_dim, hidden_dim, batch_first=True, bidirectional=True)
        self.fc = nn.Linear(hidden_dim * 2, output_dim) # Bidirectional

    def forward(self, features):

        seq_len = features.size(2) * features.size(3) # Assuming height x width gives the sequence length
        # print("Features shape before reshape:", features.shape)

        features = features.reshape(features.size(0), seq_len, -1) # Flatten height and width dimensions
        # print("Features shape after reshape:", features.shape)

        output, _ = self.rnn(features)
        # print(output.shape)
        output = self.fc(output)
        # print(output.shape)
        return output # Character probabilities

class OCRModel(nn.Module):
    def __init__(self, cnn_encoder, rnn_decoder):
        super(OCRModel, self).__init__()
        self.encoder = cnn_encoder
        self.decoder = rnn_decoder

    def forward(self, x):
        features = self.encoder(x)

```

```
        output = self.decoder(features)
        return output
```

## Training

```
def train(model, train_loader, val_loader, criterion, optimizer, num_epochs=10):
    for epoch in range(num_epochs):
        model.train() # Set model to training mode
        total_train_loss = 0
        batch_train_losses = [] # To store loss per batch
        in training

            for batch_idx, (images, labels, label_lengths) in enumerate(train_loader): # Assuming dataloader provides batches of images and labels
                optimizer.zero_grad()

                # Get the model outputs (shape: batch_size x max_seq_len x num_classes)
                outputs = model(images) # (N, T, C) where N is batch size, T is time steps, C is num classes
                outputs = outputs.permute(1, 0, 2) # (T, N, C)
                for CTCLoss (time x batch x classes)

                    # Compute the lengths of the input sequences
                    (i.e., the number of time steps)
                    input_lengths = torch.full(size=(images.size(0),), fill_value=outputs.size(0), dtype=torch.long)

                    # Flatten the labels and create target_lengths
                    (length of each label sequence)
                    flattened_labels = torch.cat([label for label in labels]) # Concatenate the label tensors
                    target_lengths = torch.tensor([len(label) for label in labels], dtype=torch.long)

                    # Compute the CTC loss
```

```

        loss = criterion(outputs, flattened_labels, input_lengths, target_lengths)

        loss.backward()
        optimizer.step()
        total_train_loss += loss.item() # Accumulate the loss for monitoring
        batch_train_losses.append(loss.item()) # Store loss per batch

        # Print the average training loss for the entire epoch
        avg_train_loss = abs(total_train_loss) / len(train_loader)
        print(f"Epoch [{epoch+1}/{num_epochs}], Average Train Loss: {avg_train_loss:.4f}")

        # Now calculate the validation loss
        model.eval() # Set model to evaluation mode
        total_val_loss = 0
        batch_val_losses = [] # To store loss per batch in validation
        with torch.no_grad(): # Disable gradients during validation
            for batch_idx, (images, labels, label_lengths) in enumerate(val_loader): # Assuming val_loader is available
                # Get the model outputs (shape: batch_size x max_seq_len x num_classes)
                outputs = model(images) # (N, T, C) where N is batch size, T is time steps, C is num classes
                outputs = outputs.permute(1, 0, 2) # (T, N, C) for CTCLoss (time x batch x classes)

                # Compute the lengths of the input sequences (i.e., the number of time steps)
                input_lengths = torch.full(size=(images.size(0),), fill_value=outputs.size(0), dtype=torch.long)

```

```

                # Flatten the labels and create target_lengths
                # (length of each label sequence)
                flattened_labels = torch.cat([label for label in labels]) # Concatenate the label tensors
                target_lengths = torch.tensor([len(label) for label in labels], dtype=torch.long)

                # Compute the CTC loss
                loss = criterion(outputs, flattened_labels,
input_lengths, target_lengths)
                total_val_loss += loss.item() # Accumulate
the validation loss
                batch_val_losses.append(loss.item()) # Store loss per batch

                # Print the average validation loss for the entire
epoch
                avg_val_loss = abs(total_val_loss) / len(val_loader)
            print(f"Epoch [{epoch+1}/{num_epochs}], Average Validation Loss: {avg_val_loss:.4f}")

```

```

Epoch [1/10], Average Train Loss: 25.7310
Epoch [1/10], Average Validation Loss: 24.5117
Epoch [2/10], Average Train Loss: 20.5848
Epoch [2/10], Average Validation Loss: 22.0605
Epoch [3/10], Average Train Loss: 16.4678
Epoch [3/10], Average Validation Loss: 19.8545
Epoch [4/10], Average Train Loss: 13.1743
Epoch [4/10], Average Validation Loss: 17.8690
Epoch [5/10], Average Train Loss: 10.5394
Epoch [5/10], Average Validation Loss: 16.0821
Epoch [6/10], Average Train Loss: 8.4315
Epoch [6/10], Average Validation Loss: 14.4739
Epoch [7/10], Average Train Loss: 6.7452
Epoch [7/10], Average Validation Loss: 13.0265
Epoch [8/10], Average Train Loss: 5.3962

```

```
Epoch [8/10], Average Validation Loss: 11.7239
Epoch [9/10], Average Train Loss: 4.3169
Epoch [9/10], Average Validation Loss: 10.5515
Epoch [10/10], Average Train Loss: 3.4536
Epoch [10/10], Average Validation Loss: 9.4963
```

## Evaluation

Test loader ANCC = 0.6142

Random baseline ANCC = 0.0278

Decoded predictions and ground truth

Prediction: abasuse

abature

Prediction: abassareise

abastardize